

---

# CS2030 Lecture 10

## From Sequential to Parallel Streams

Henry Chia (hchia@comp.nus.edu.sg)

Semester 1 2019 / 2020

# Lecture Outline

---


- Concurrency versus parallelism
- Parallel streams
  - Debugging parallel streams
  - Comparing sequential and parallel streams
- Correctness of parallel streams
  - reduce operator
  - Accumulator and combiner
- Fork and join in parallel streams
  - Overhead of parallelization

# Concurrency vs Parallelism

---

- A single core processor executes one instruction at a time
  - Only one process can run at any one time
  - Context-switching allows multi-tasking on a single processor
- Concurrent programs run concurrently via threads
  - OS switches between threads
  - Separate unrelated tasks into separate threads
  - Improves processor utilization
- Parallel computing involves multiple subtasks running at the same time on multiple (possibly multi-core) processors
- Parallel programs are concurrent, but not all concurrent programs are parallel

# Concurrency vs Parallelism



# Parallel Streams

---

- Parallel streams use a common `ForkJoinPool` via the static `ForkJoinPool.commonPool()` method

```
jshell> ForkJoinPool.commonPool().getParallelism()  
$1 ==> 23
```

- The level of parallelism can be controlled by setting the system property during program run

```
System.setProperty("java.util.concurrent.ForkJoinPool.common.parallelism", "4")
```

or including the following flag when running the program

```
$ java -Djava.util.concurrent.ForkJoinPool.common.parallelism=4 ...
```

- Similar to the method `stream()`, Java `Collection(s)` also support the method `parallelStream()` to create a parallel stream of elements


# Parallel Streams

- Using prime number testing as an example

```
static boolean isPrime(int n) {  
    return IntStream  
        .rangeClosed(2, (int) Math.sqrt(n))  
        .noneMatch(x -> n % x == 0);  
}
```

- Count number of primes between 2,000,000 and 3,000,000

```
long count = IntStream.range(2_000_000, 3_000_000)  
    .filter(x -> isPrime(x))  
    .count();
```




# Parallel Streams

- Parallelizing the search for primes

```
long count = IntStream.range(2_000_000, 3_000_000)
    .parallel()
    .filter(x -> isPrime(x))
    .count();
```

- The `parallel()` intermediate operation turns on a boolean flag that switches the stream pipeline to be parallel
  - Invoked anywhere between the data source and terminal
  - The counter operation is `sequential()`



# Parallel Streams

---

```
public static void main(String[] args) {  
    if (args.length != 0) {  
        System.setProperty(  
            "java.util.concurrent.ForkJoinPool.common.parallelism",  
            args[0]);  
    }  
    System.out.println("Number of worker threads: " +  
        ForkJoinPool.commonPool().getParallelism());  
  
    Instant start = Instant.now();  
    long howMany = IntStream.range(2_000_000, 3_000_000)  
        .parallel()  
        .filter(x -> isPrime(x))  
        .count();  
    Instant stop = Instant.now();  
  
    System.out.println(howMany + " : " +  
        Duration.between(start, stop).toMillis() + "ms");  
}
```



# Debugging Parallel Streams

---

- To time the execution of a process,
  - `java.time.Instant`'s `now()` method returns the current `Instant` from the system clock
  - `java.time.Duration`'s `between()` returns the `Duration` of two `Instances` (an implementation of `Temporal`)
  - `Duration`'s `toMillis()/toNanos()/...` extracts the desired representation of the duration

```
java.util.Instant;  
java.util.Duration;
```

```
Instant start, stop;  
start = Instant.now();  
/* perform some task */  
stop = Instant.now();
```

```
long timeInMillis = Duration.between(start, stop).toMillis();
```

# Debugging Parallel Streams

---

- To debug and manage each execution thread
  - `Thread.currentThread()` (or `Thread.currentThread().getName()`) to retrieve the identity of the thread
  - `Thread.sleep(long millis)` causes the currently executing thread to sleep (i.e. temporarily cease execution) for the specified number of milliseconds
    - ▷ Used within a **try.. catch** block
    - ▷ Example, letting a thread sleep for one second

```
try {  
    ...  
    Thread.sleep(1000);  
    ...  
} catch (InterruptedException e) { }
```

# Debugging Parallel Streams

- Effect of parallelizing a stream

```
int sum = IntStream.of(1, 2, 3, 4, 5)
    .parallel()
    .filter(x -> {
        System.out.println("filter: " + x + " "
            + Thread.currentThread().getName());
        return x % 2 == 1;
    })
    .map(x -> {
        System.out.println("map: " + x + " "
            + Thread.currentThread().getName());
        return x;
    })
    .reduce(0, (x, y) -> {
        System.out.println("reduce: " + x + " + " + y + " "
            + Thread.currentThread().getName());
        return x + y;
    });
System.out.println(sum);
```

# Comparing Sequential and Parallel Streams

- Suppose given the following task unit

```
class OneSecondTask {  
    int ID;  
  
    public OneSecondTask(int ID) {  
        this.ID = ID;  
    }  
  
    public int compute() {  
        System.out.println(Thread.currentThread().getName());  
        try {  
            Thread.sleep(1000);  
        } catch (InterruptedException e) {  
            throw new RuntimeException(e);  
        }  
        return ID;  
    }  
}  
//      completableFutureThenAccept();
```

- Comparing sequential and parallel streams:

```
Stream<OneSecondTask> tasks = IntStream.range(0, 10)  
    .mapToObj(x -> new OneSecondTask(x));  
runTasks(tasks);
```

# Comparing Sequential and Parallel Streams

```
public static void runTasks(Stream<OneSecondTask> tasks) {  
    Instant start = Instant.now();  
    List<Integer> result = tasks  
        .map(x -> x.compute())  
        .collect(Collectors.toList());  
    Instant stop = Instant.now();  
    System.out.print(result + " ");  
    System.out.println(Duration.between(start, stop).toMillis() + "ms");  
}
```

- Sequential stream on 4 worker threads:

```
main  
main  
main  
main  
main  
main  
main  
main  
main  
main  
main  
[0, 1, 2, 3, 4, 5, 6, 7, 8, 9] 10003ms
```

# Comparing Sequential and Parallel Streams

---

```
Stream<OneSecondTask> tasks = IntStream.range(0, 10)
    .mapToObj(x -> new OneSecondTask(x));
runTasks(tasks.parallel());
```

- Parallel stream on 4 worker threads:

```
main
ForkJoinPool.commonPool-worker-1
ForkJoinPool.commonPool-worker-3
ForkJoinPool.commonPool-worker-2
ForkJoinPool.commonPool-worker-4
ForkJoinPool.commonPool-worker-3
ForkJoinPool.commonPool-worker-2
ForkJoinPool.commonPool-worker-4
ForkJoinPool.commonPool-worker-1
ForkJoinPool.commonPool-worker-3
[0, 1, 2, 3, 4, 5, 6, 7, 8, 9] 3006ms
```

# Correctness of Parallel Streams

- To ensure correct parallel execution, stream operations
  - must not interfere with stream data (true for sequential streams also)
  - preferably stateless with no side effects
- Example of interference:

For parallel streams to work properly, there should be ideally no sideeffects and immutability

```
List<String> list = new ArrayList<>(
    List.of("abc", "def", "xyz"));

list.stream()
    .peek(str -> {
        if (str.equals("xyz")) {
            list.add("pqr");
        }
    })
    .forEach(x -> {});
```

# Correctness of Parallel Streams

- Example of side-effects:

```
List<Integer> list = new ArrayList<>(
    Arrays.asList(1, 3, 5, 7, 9, 11, 13, 15, 17, 19));
List<Integer> result = new ArrayList<>();
```

- The following is erroneous

```
list.parallelStream()
    .filter(x -> isPrime(x))
    .forEach(x -> result.add(x));
```

- Use `.collect` instead

```
result = list.parallelStream()
    .filter(x -> isPrime(x))
    .collect(Collectors.toList());
```

- Can also consider using `forEachOrdered()`, or a thread-safe `list CopyOnWriteArrayList`

in order

not in order



# Inherently Parallelizable reduce

- Consider Stream's three-argument reduce method:  

```
<U> U reduce(U identity,  
             BiFunction<U,? super T,U> accumulator,  
             BinaryOperator<U> combiner)
```
- Rules to follow when parallelizing
  - `combiner.apply(identity, i)` must be equal to `i`
  - `combiner` and `accumulator` must be associative, i.e. order of application does not matter
  - `combiner` and `accumulator` must be compatible, i.e.  
`combiner.apply(u, accumulator.apply(identity, t))`  
must be equal to `accumulator.apply(u, t)`

# Accumulator and Combiner


- The following example compiles with the above rules:

```
Stream.of(1,2,3,4)
    .parallel()
    .reduce(1, (x,y) -> x + y, (x,y) -> x + y)
```

- To see the effects of accumulator and combiner:

```
sum = Stream.of(1, 2, 3, 4, 5)
    .parallel()
    .filter(x -> {
        System.out.println("filter: " + x + " "
            + Thread.currentThread().getName());
        return x % 2 == 1;
    })
    .reduce(0,
        (x, y) -> {
            System.out.println("accumulate: " + x + " + " + y + " "
                + Thread.currentThread().getName());
            return x + y;
        },
        (x, y) -> {
            System.out.println("combine: " + x + " + " + y + " "
                + Thread.currentThread().getName());
            return x + y;
        }
    );
```

# Accumulator and Combiner



# Accumulator and Combiner

- Output from a sample run:

```
filter: 5 ForkJoinPool.commonPool-worker-1
filter: 4 ForkJoinPool.commonPool-worker-3
filter: 1 ForkJoinPool.commonPool-worker-3
filter: 3 main
filter: 2 ForkJoinPool.commonPool-worker-2
accumulate: 0 + 5 ForkJoinPool.commonPool-worker-1 // (A)
accumulate: 0 + 1 ForkJoinPool.commonPool-worker-3
combine: 1 + 0 ForkJoinPool.commonPool-worker-3 // (C)
accumulate: 0 + 3 main
combine: 0 + 5 ForkJoinPool.commonPool-worker-1
combine: 3 + 5 ForkJoinPool.commonPool-worker-1 // (B)
combine: 1 + 8 ForkJoinPool.commonPool-worker-1
9
```


- (A) Accumulation with identity
- (B) Combining results from accumulation
- (C) Combining with identity

# Accumulator and Combiner

- Erroneous examples where rules are violated
  - `combiner.apply(identity, i)` not equal to `i`  
**double** `result = Stream`  
    `.of(1, 2, 3, 4)`  
    `.parallel()`  
    `.reduce(1.0,`  
        `(x, y) -> x + y,`  
        `(x, y) -> x + y);`
  - for division, the order of application **does** matter

```
double result = Stream
    .of(1, 2, 3, 4)
    .parallel()
    .reduce(24.0,
        (x, y) -> 1.0 * x / y,
        (x, y) -> 1.0 * x / y);
```

# Fork and Join in Parallel Streams



- `parallel()` runs `fork` to create sub-tasks running the same chain of operations on sub-streams
  - Processes for sub-tasks are run in multiple threads when appropriate
  - Threads are shared from a common **Fork Join Pool**
- `combiner` in `reduce` runs `join` to combine the results

# Fork and Join in Parallel Streams

- Should we exploit parallelism to the fullest?

```
return IntStream
    .rangeClosed(2, (int) Math.sqrt(n))
    .parallel()
    .noneMatch(x -> n % x == 0);
```

- Parallelizing a trivial task actually creates more work in terms of parallelizing overhead
- Parallelization is worthwhile only if the task is complex enough that the benefit of parallelization outweighs the overhead
  - In primality testing, checking  $(n \% x == 0)$  is trivial;
  - Parallelizing it induces more overhead in terms of processing the forks and joins
- Holds true for all parallel and concurrent programs

# Lecture Summary

---

- Familiarity with the use of sequential and parallel streams
- Able to compare performances between sequential and parallel streams
- Able to debug parallel streams
- Adherence to rules for parallelizing streams
- Appreciate fork and join in parallel streams
- Appreciate fork/join overhead