

---

# CS2030 Lecture 11

## Fork/Join Framework

Henry Chia (hchia@comp.nus.edu.sg)

Semester 1 2019 / 2020

# Lecture Outline

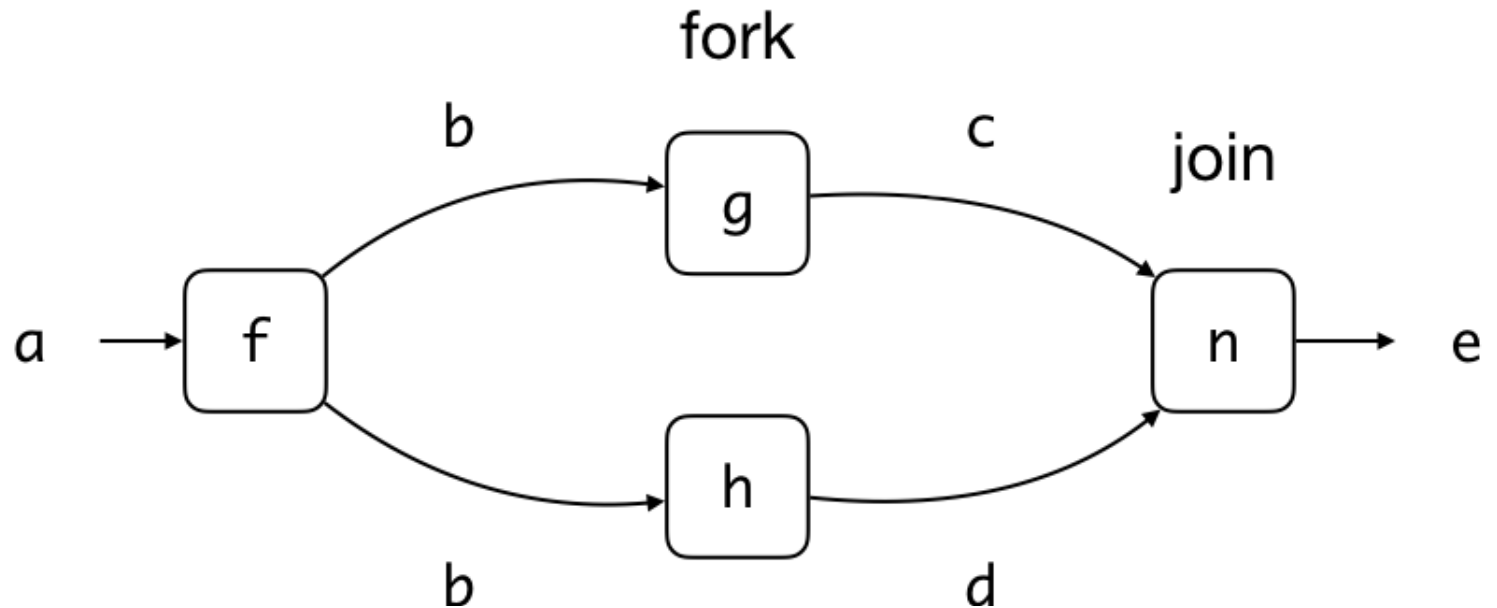
---

- Fork and join tasks
- Java's Fork/join framework
  - Sub-classing a `RecursiveTask`
- Thread pools
  - Global queue
  - Local deque (double-ended queue)
- Work stealing
- Order of fork and join
- Overhead of fork and join

# Fork and Join

- Given the following program fragment and *computation* graph

```
b = f(a);  
c = g(b);  
d = h(b);  
e = n(c,d);
```



- $f(a)$  invoked before  $g(b)$  and  $h(b)$ ;  $n(c,d)$  invoked after
- How about the order of  $g(b)$  and  $h(b)$ ?
  - If  $g$  and  $h$  does not produce side effects, then parallelize
  - **Fork** task  $g$  to execute at the same time as  $h$ , and **join** back task  $g$  later

# Example: Summing an Array... *Recursively*

```
class Summer {  
    static int threshold;  
  
    static int sumLeftRight(int[] array, int low, int high) {  
        if (high - low < threshold) {  
            int sum = 0;  
            for (int i = low; i <= high; i++) {  
                sum += array[i];  
            }  
            return sum;  
        } else {  
            Similar to the mergesort algorithm  
=> Divide and Conquer Algo  
            int middle = (low + high) / 2;  
            int leftSum = sumLeftRight(array, low, middle);  
            int rightSum = sumLeftRight(array, middle + 1, high);  
            return leftSum + rightSum;  
        }  
    }  
}  
  
int[] array = IntStream.rangeClosed(1, 10).toArray();  
Summer.threshold = Integer.parseInt(args[0]);  
int sum = sumLeftRight(array, 0, array.length - 1);
```

# Transform to a Recursive Class

```
class Summer {
    private int[] array;
    private int low;
    private int high;
    static int threshold;

    Summer(int[] array, int low, int high) {
        this.array = array;
        this.low = low;
        this.high = high;
    }

    int compute() {
        if (high - low < threshold) {
            int sum = 0;
            for (int i = low; i <= high; i++) {
                sum += array[i];
            }
            return sum;
        } else {
            int middle = (low + high) / 2;
            Summer left = new Summer(array, low, middle);
            Summer right = new Summer(array, middle + 1, high);
            return left.compute() + right.compute();
        }
    }
}
```

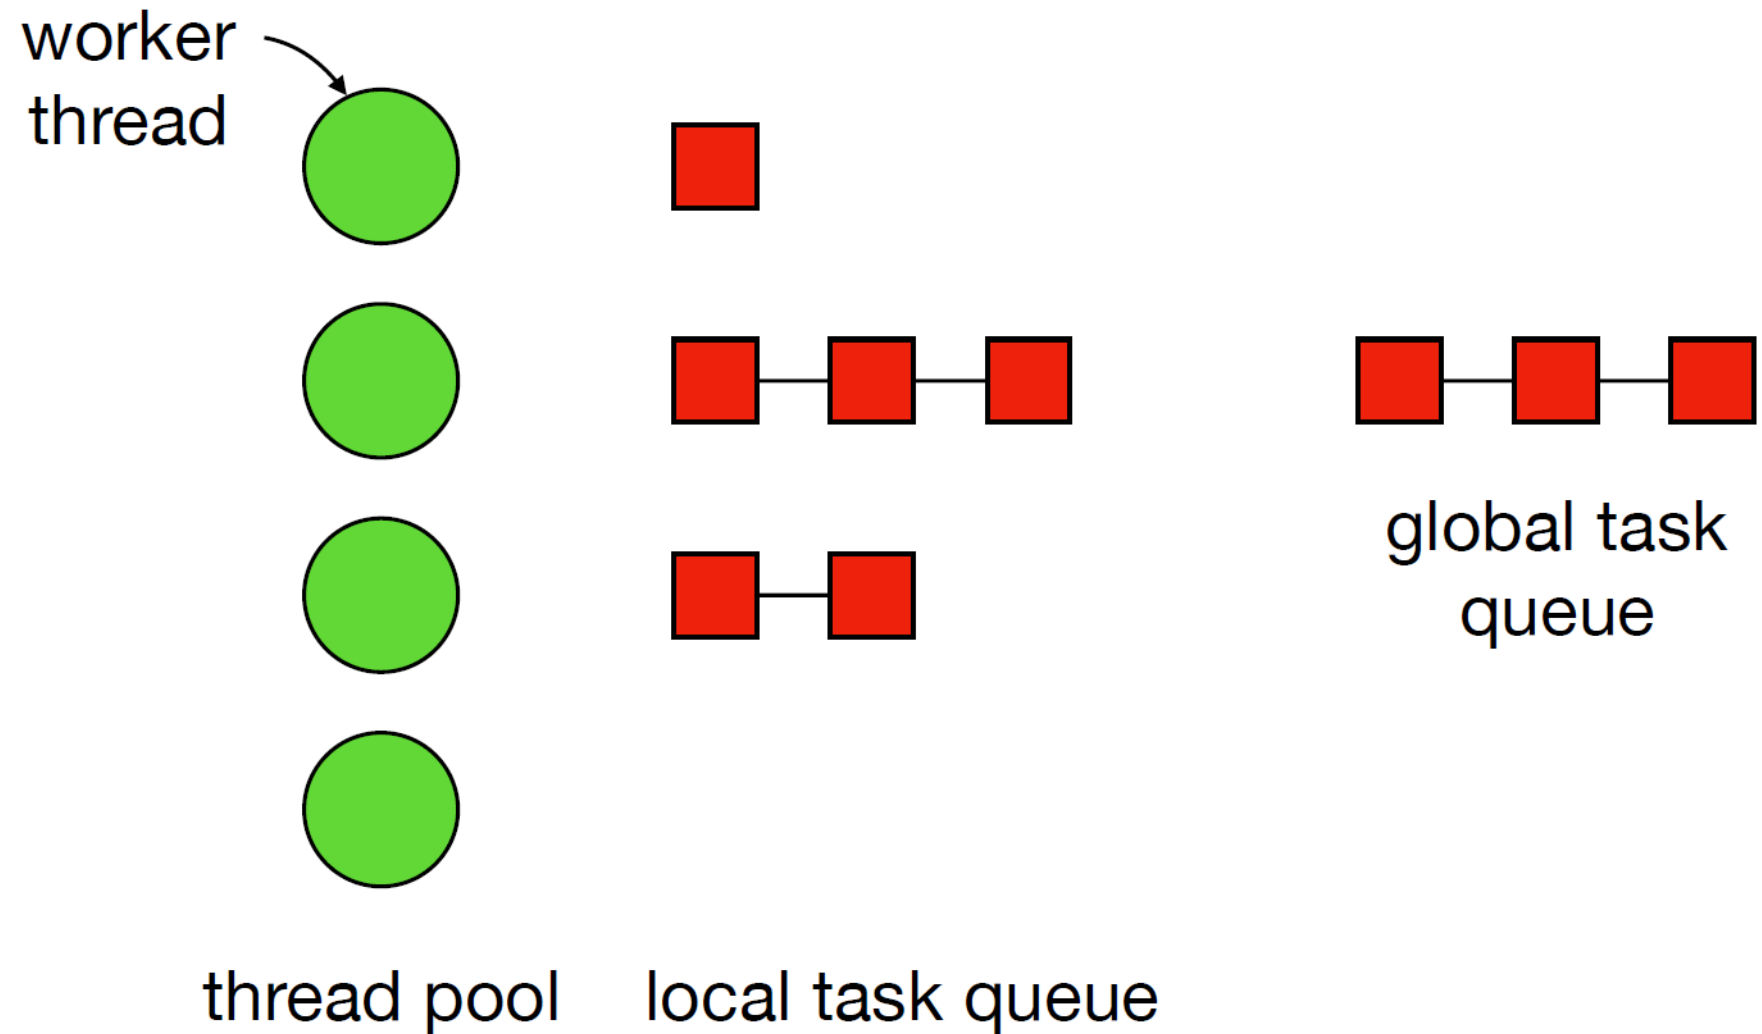
# Subclassing RecursiveTask<T> for Fork/Join

```
import java.util.concurrent.RecursiveTask;

class Summer extends RecursiveTask<Integer> {
    ...
    @Override
    protected Integer compute() {
        if (high - low < threshold) {
            int sum = 0;
            for (int i = low; i < high; i++) {
                sum += array[i];
            }
            return sum;
        } else {
            int middle = (low + high) / 2;
            Summer left = new Summer(low, middle, array);
            Summer right = new Summer(middle, high, array);

            left.fork();
            return right.compute() + left.join();
        }
    }
}
```

# Thread Pools



# Thread Pools

---

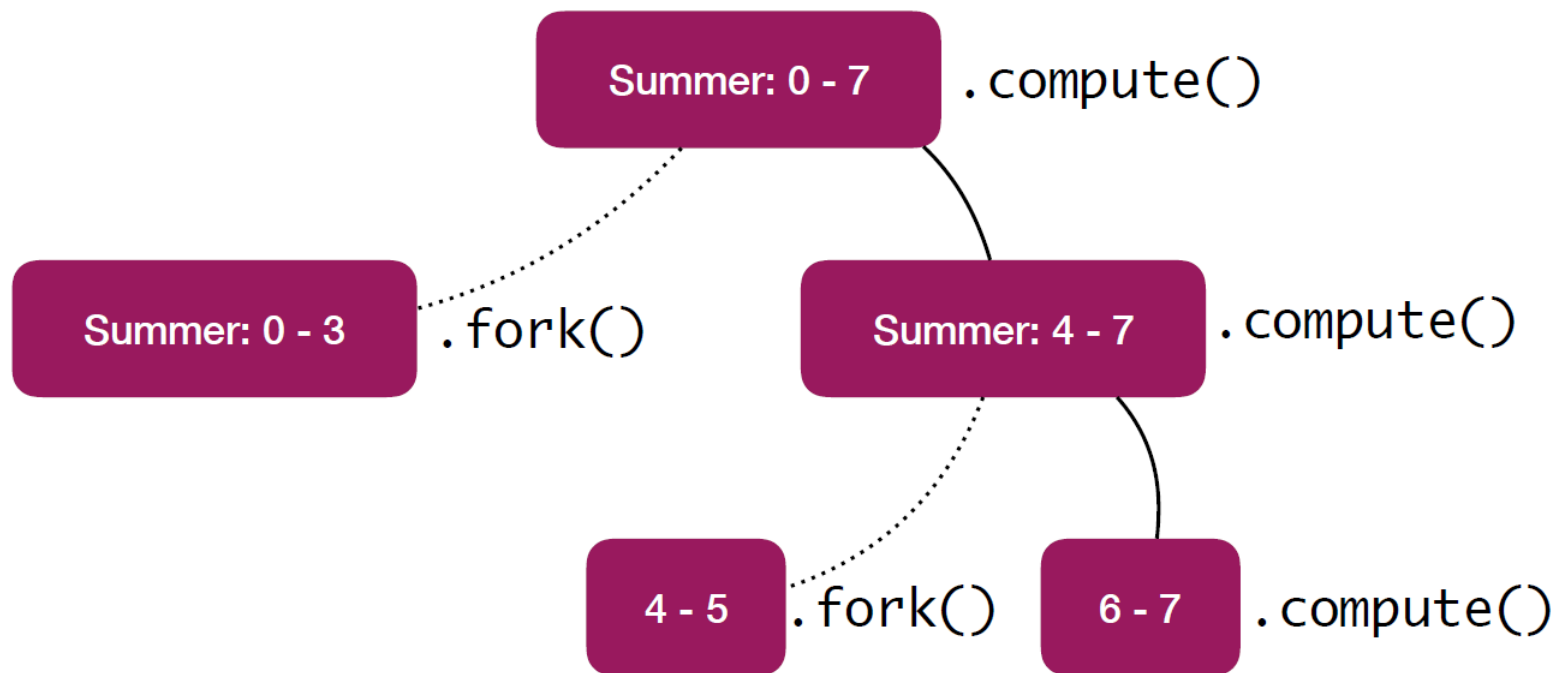
- Java maintains a pool of *worker threads*
  - Each thread is an abstraction of a running task
  - Task submitted to the pool for execution, and joins the global queue or worker queue
  - Worker thread picks a task from the queue to execute
- `ForkJoinPool` is the class that implements the thread pool for `RecursiveTask` (a sub-class of `ForkJoinTask`)
- To submit task to the thread pool for execution, either
  - `task.compute()` that invokes task immediately; may result in stack overflow if too many recursive tasks
  - `invoke(task)` that gets the task to join the queue, waiting to be carried out by a worker (recommended)



# Example: Summing an Array

- Summing array of eight elements with threshold set to 2

0	1	2	3	4	5	6	7
3	3	1	3	6	8	4	2



# Example: Summing an Array

```
@Override
protected Integer compute() {
    System.out.println(low + "," + high + ":" + Thread.currentThread().getName());
    if (high - low < threshold) {
        int sum = 0;
        for (int i = low; i < high; i++) {
            sum += array[i];
        }
        return sum;
    } else {
        int middle = (low + high) / 2;
        Summer left = new Summer(array, low, middle);
        Summer right = new Summer(array, middle + 1, high);
        left.fork();
        return right.compute() + left.join();
    }
}
```

□ Running with ForkJoinPool.common.parallelism=2


```
0,7:main
4,7:main
6,7:main
0,3:ForkJoinPool.commonPool-worker-1
4,5:ForkJoinPool.commonPool-worker-2
2,3:ForkJoinPool.commonPool-worker-1
0,1:ForkJoinPool.commonPool-worker-2
```

refer to the diagram next page


```
0,7:main
4,7:main
6,7:main
4,5:main
0,3:ForkJoinPool.commonPool-worker-2
0,1:ForkJoinPool.commonPool-worker-1
2,3:ForkJoinPool.commonPool-worker-2
```

# Queuing of Forked Tasks

- main thread performs compute on task  $[0, 7]$




- main forks task  $[0, 3]$  to global queue, then computes  $[4, 7]$




# Work Fetching from Global Queue

- Worker #1 fetches task [0-3] from global queue




- main forks [4-5] & computes [6-7]; worker #2 fetches [4-5]




# Work Stealing

- Worker #1 forks [0-1] and computes [2-3]




- Worker #2 completes [4-5] returns, and steals [0-1]



# Local Deque (Double-Ended Queue)

- E.g., invoking `ForkJoinPool.commonPool().invoke(task)`
- Worker #1 processes 0-7: forks [0-3] and computes [4-7]
- Worker #1 processes 4-7: forks [4-5] and computes [6-7]




- When worker #1 completes [6-7], it makes more sense for him to work on [4-5], so as to complete [4-7]
- Local task queue is a double ended queue
  - Forked tasks added to the **head** of the queue
  - Steal tasks from the **end of the queue**
  - Rational: bigger tasks are stolen; smaller ones self-served


So that 4-7 is completed and the size of the stack is decreased

# Work Stealing from Back of Dequeue

- Worker #3 steals task [0-3]




- Worker #3 forks [0-1] (worker #2 steals) and computes [2-3]




# Another Possible Scenario...

- Worker #1 completes [6-7], but worker #2 has not



- Worker #3 completes [2-3], but is now blocked waiting for worker #2 to return with a value
- Worker #1 can then service [4-5] from the head of its queue





# Order of Fork/Join

```
@Override
protected Integer compute() {
    System.out.println(low + ", " + high + ":" + Thread.currentThread().getName());
    if (high - low < threshold) {
        int sum = 0;
        for (int i = low; i < high; i++) {
            sum += array[i];
        }
        return sum;
    } else {
        int middle = (low + high) / 2;
        Summer left = new Summer(array, low, middle);
        Summer right = new Summer(array, middle + 1, high);

        left.fork();
        right.fork();

        return right.join() + left.join();
    }
}
```

Main does 0,7  
fork left(4,7) fork right (0,3)  
now the queue: (4,7) (0,3)  
self service from the front can only pick up from the front but task (4,7) is blocking task (0,3)

- ☐ How about using only forks and joins
- ☐ Does the ordering matter?

# Order of Fork/Join

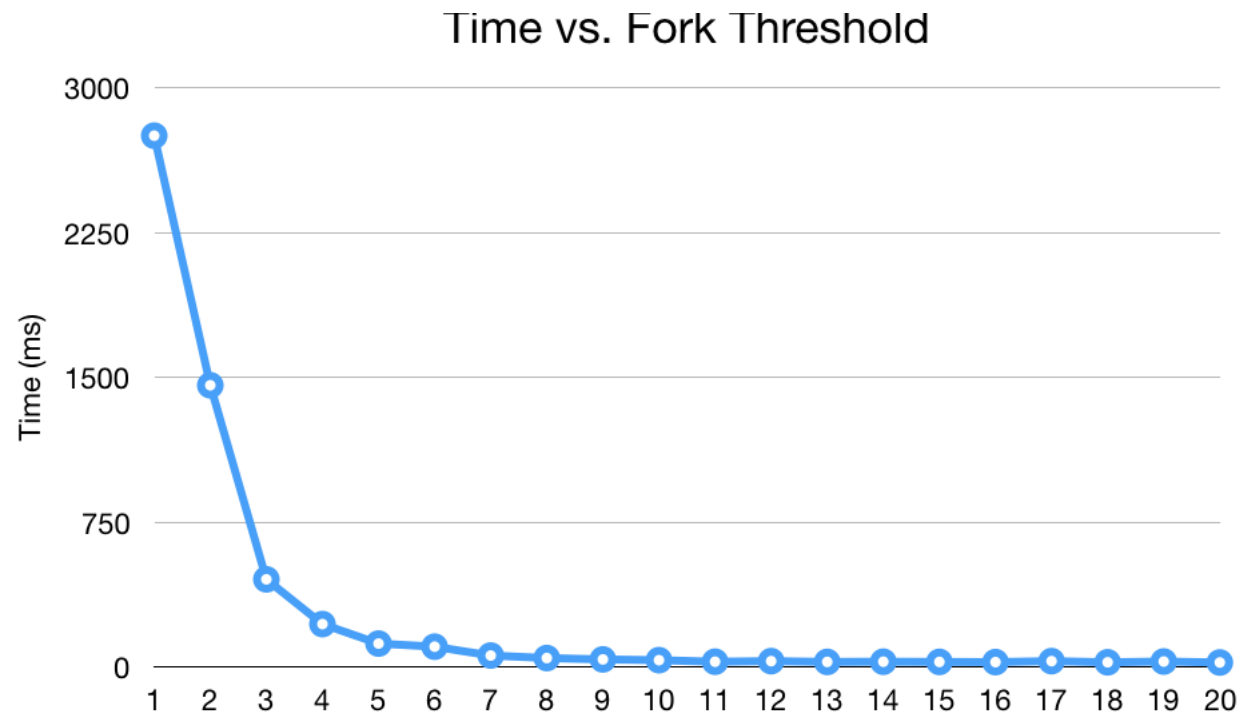
- Fork-join pair acts like a call (fork) and return (join) from a parallel recursive function.
  - Returns (joins) should be performed innermost-first
  - Performing

```
a.fork(); b.fork();  
b.join(); a.join();
```

FIFO  
First In First Out  
The first Task added must be computed first
- is likely to be substantially more efficient than joining task a before task b
- Work-stealing threadpools have a fixed number of threads
  - Any blocking operation in one of these threads will reduce overall performance

# Overhead of Fork/Join

- Forking and joining creates additional overhead
  - wrap the computation in an object
  - submit object to a queue of tasks
  - workers go through the queue to execute tasks



# Lecture Summary

---

- Appreciate the use of fork and join in parallel/concurrent programming
- Understand how tasks are forked in a local deque, and why the ordering of forks and joins matter
- Understand how work stealing distributes tasks among worker threads
- Appreciate the overhead involved in parallelizing using fork and join